

CaseStudy_EDA

Import Libraries and Dataset

```
In [23]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from google.colab import drive
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from scipy.stats import chi2_contingency, ttest_ind

# Mount Drive and Load Dataset
drive.mount('/content/drive')
train_data = '/content/drive/MyDrive/AIAD_Casestudy/train.csv'
df = pd.read_csv(train_data)
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Data Visualisation

```
In [24]: df.head()
```

Out[24]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450

In [25]:

```
# Data Overview
df.info()
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null    int64
1   Survived        891 non-null    int64
2   Pclass          891 non-null    int64
3   Name            891 non-null    object
4   Sex             891 non-null    object
5   Age            714 non-null    float64
6   SibSp           891 non-null    int64
7   Parch           891 non-null    int64
8   Ticket          891 non-null    object
9   Fare            891 non-null    float64
10  Cabin           204 non-null    object
11  Embarked        889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

Out[25]:

	PassengerId	Survived	Pclass	Age	SibSp	Parch
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000

General overview of the dataset, take into consideration that there are null values in "age", "cabin" and "embarked"

In [26]:


```
# Missing Values
missing_data = df.isnull().sum().sort_values(ascending=False)
missing_percentage = (df.isnull().sum() / len(df)) * 100
missing_data_df = pd.DataFrame({'Missing Values': missing_data, 'Percentage': missing_percentage})
print(missing_data_df)
```

	Missing Values	Percentage
Age	177	19.865320
Cabin	687	77.104377
Embarked	2	0.224467
Fare	0	0.000000
Name	0	0.000000
Parch	0	0.000000
PassengerId	0	0.000000
Pclass	0	0.000000
Sex	0	0.000000
SibSp	0	0.000000
Survived	0	0.000000
Ticket	0	0.000000

This confirms null values within the dataset

In [27]:

```
# Visualize Missing Values
plt.figure(figsize=(10, 6))
sns.heatmap(df.isnull(), cbar=False, cmap='viridis')
plt.title('Missing Values Heatmap')
plt.show()
```



Visualising Null Values

```
In [28]: # Fill Missing Values
df['Age'] = df['Age'].fillna(df.groupby(['Pclass', 'Sex'])['Age'].transform('median'))
df['Embarked'].fillna(df['Embarked'].mode()[0], inplace=True)
```

<ipython-input-28-9666381a2d49>:3: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Embarked'].fillna(df['Embarked'].mode()[0], inplace=True)
```

To fill up missing values in "age", i decided to use the median of pclass and sex, for "embarked" i used the most common port.


```
In [29]: # Missing Values
missing_data = df.isnull().sum().sort_values(ascending=False)
missing_percentage = (df.isnull().sum() / len(df)) * 100
missing_data_df = pd.DataFrame({'Missing Values': missing_data, 'Percentage': missing_percentage})
```

```
print(missing_data_df)
```

	Missing Values	Percentage
Age	0	0.000000
Cabin	687	77.104377
Embarked	0	0.000000
Fare	0	0.000000
Name	0	0.000000
Parch	0	0.000000
PassengerId	0	0.000000
Pclass	0	0.000000
Sex	0	0.000000
SibSp	0	0.000000
Survived	0	0.000000
Ticket	0	0.000000

this shows that there are no more nulls except for cabin, i decided to use cabin as a feature

```
In [30]: # Visualize Missing Values
plt.figure(figsize=(10, 6))
sns.heatmap(df.isnull(), cbar=False, cmap='viridis')
plt.title('Missing Values Heatmap')
plt.show()
```




Feature Visualisation

```
In [31]: f, ax = plt.subplots(1, 2, figsize=(18, 8))

df['Survived'].value_counts().plot.pie(
    explode=[0, 0.1],
    autopct='%1.1f%%',
    ax=ax[0],
)
ax[0].set_title('Survived')
ax[0].set_ylabel('')

sns.countplot(x='Survived', data=df, ax=ax[1])
ax[1].set_title('Survived')

plt.show()
```



As seen, more than 60% of people that were on the Titanic perished, not so unsinkable.

Feature Engineering

```
In [33]: # New Features
df['HasCabin'] = df['Cabin'].notna().astype(int)
df['FamilySize'] = df['SibSp'] + df['Parch'] + 1
df['Title'] = df['Name'].apply(lambda x: x.split(',')[1].split('.')[0].strip()
    "Mr": "Mr", "Miss": "Miss", "Mrs": "Mrs", "Master": "Master",
    "Dr": "Other", "Rev": "Other", "Col": "Other", "Mlle": "Miss",
    "Major": "Other", "Mme": "Mrs", "Capt": "Other", "Countess": "Other",
    "Jonkheer": "Other", "Don": "Other"
})
df['IsAlone'] = (df['FamilySize'] == 1).astype(int)
df['FareBin'] = pd.qcut(df['Fare'], 4, labels=[1, 2, 3, 4])

# Age Bins
```

```
df['Age_band'] = pd.cut(df['Age'], bins=[0, 16, 32, 48, 64, 80], labels=[
print(df.head())
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp
0	Braund, Mr. Owen Harris	male	22.0	1
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1
2	Heikkinen, Miss. Laina	female	26.0	0
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1
4	Allen, Mr. William Henry	male	35.0	0

	Parch	Ticket	Fare	Cabin	Embarked	HasCabin	FamilySize
0	0	A/5 21171	7.2500	NaN	S	0	2
1	0	PC 17599	71.2833	C85	C	1	2
2	0	STON/O2. 3101282	7.9250	NaN	S	0	1
3	0	113803	53.1000	C123	S	1	2
4	0	373450	8.0500	NaN	S	0	1

	Title	IsAlone	FareBin	Age_band
0	Mr	0	1	1
1	Mrs	0	4	2
2	Miss	1	2	1
3	Mrs	0	4	2
4	Mr	1	2	2

Here, i made it so that if the passenger had a cabin it could be represented by a 1 whilst if they did not it would be 0

I also combined SibSp and Parch to make Family Size

I referenced a dictionary showing the titles of people and decided to categorize them for simplicity

i also introduced a single rider feature

i also introduced a fare bin, this is to reduce the complexity of the data

age band is also a way of reducing data complexity, we can observe more categorial insights from this

```
In [34]: df.head()
```

Out[34]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450

Feature Analysis

```
In [59]: fig, axes = plt.subplots(2, 2, figsize=(14, 12))

sns.histplot(df['Age'].dropna(), kde=True, ax=axes[0, 0])
axes[0, 0].set_title("Age Distribution")

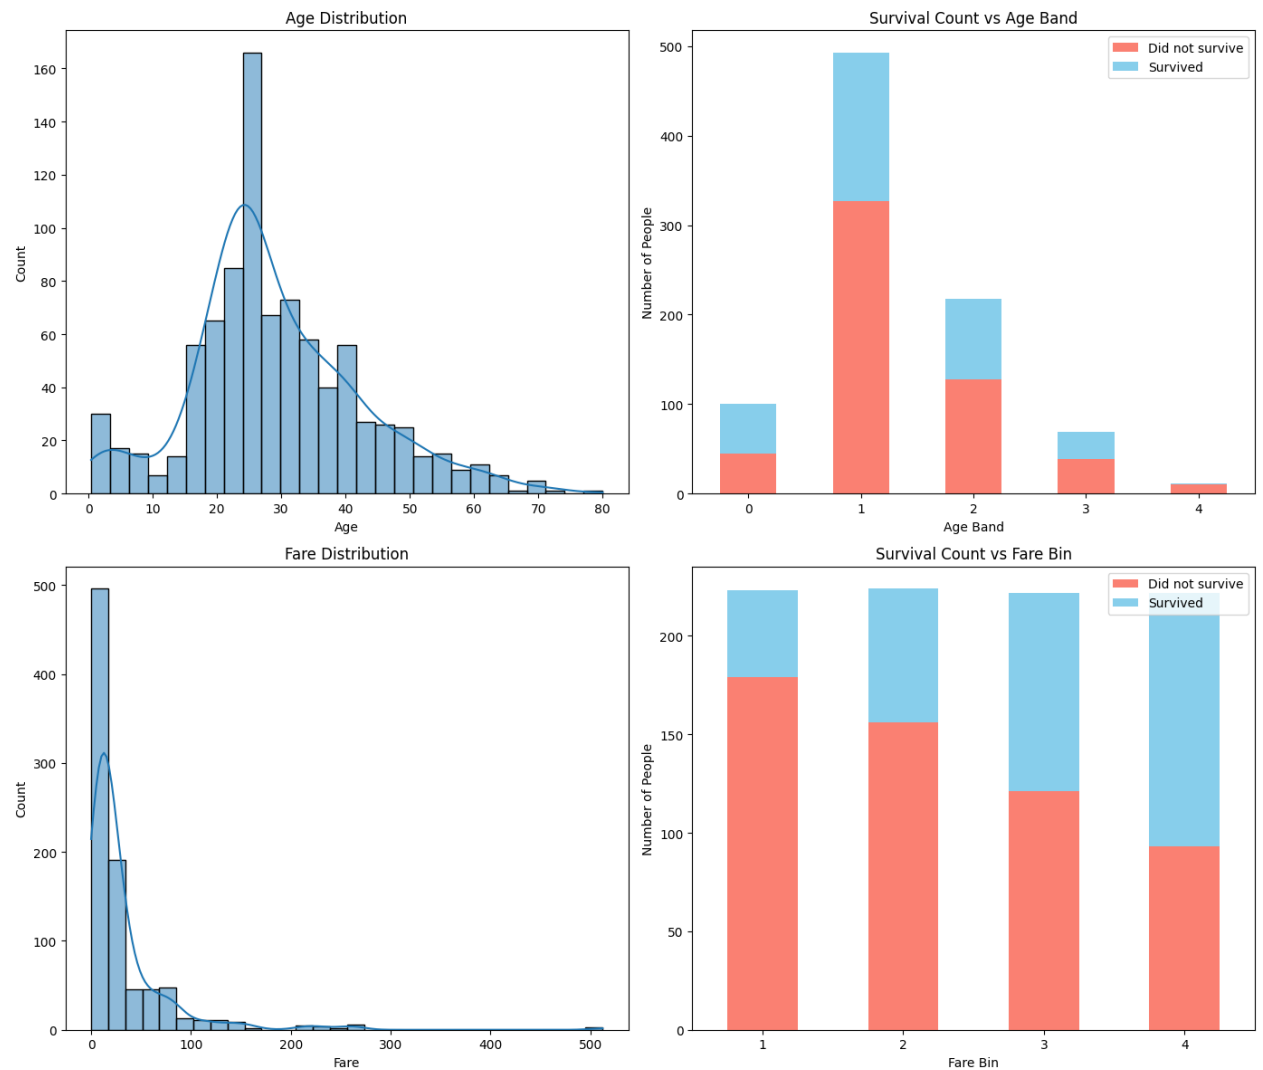
fare_survival_count_age.plot(kind='bar', stacked=True, ax=axes[0, 1], color=
axes[0, 1].set_title('Survival Count vs Age Band')
axes[0, 1].set_xlabel('Age Band')
axes[0, 1].set_ylabel('Number of People')
axes[0, 1].set_xticklabels(fare_survival_count_age.index.astype(str), rotation=
axes[0, 1].legend(['Did not survive', 'Survived'])

sns.histplot(df['Fare'], kde=True, bins=30, ax=axes[1, 0])
axes[1, 0].set_title("Fare Distribution")

fare_survival_count_fare.plot(kind='bar', stacked=True, ax=axes[1, 1], color=
axes[1, 1].set_title('Survival Count vs Fare Bin')
axes[1, 1].set_xlabel('Fare Bin')
axes[1, 1].set_ylabel('Number of People')
axes[1, 1].set_xticklabels(fare_survival_count_fare.index.astype(str), rotation=
axes[1, 1].legend(['Did not survive', 'Survived'])
```




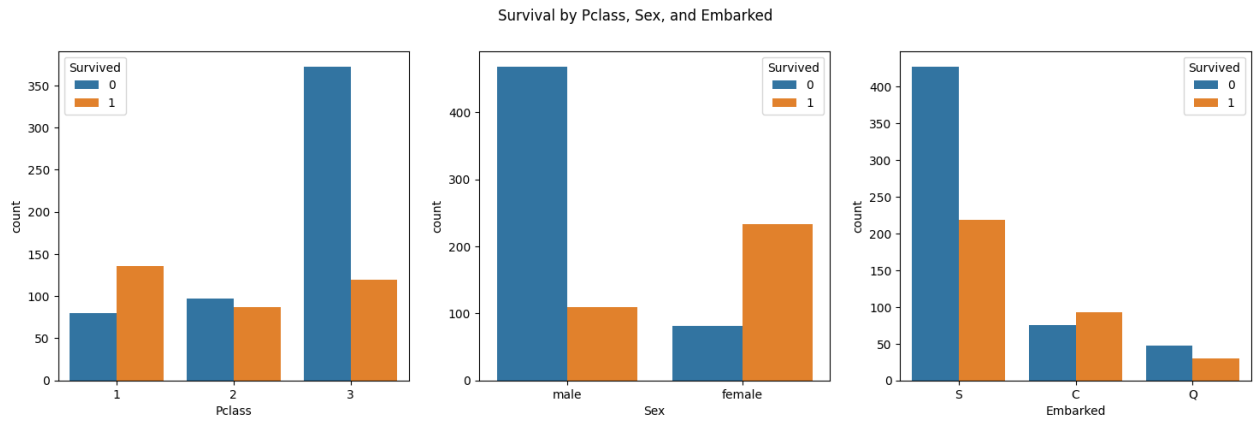
```
plt.tight_layout()
plt.show()
```



From here, we can see the general distribution for age and fare according to their hard numbers and also the categorical numbers

```
In [36]: # Survival by Categorical Features
fig, axes = plt.subplots(1, 3, figsize=(18, 5))
sns.countplot(data=df, x='Pclass', hue='Survived', ax=axes[0])
sns.countplot(data=df, x='Sex', hue='Survived', ax=axes[1])
sns.countplot(data=df, x='Embarked', hue='Survived', ax=axes[2])
fig.suptitle('Survival by Pclass, Sex, and Embarked')
plt.show()

# KDE Plot: Age Distribution
sns.kdeplot(data=df[df['Survived'] == 1], x='Age', fill=True, label='Surv')
sns.kdeplot(data=df[df['Survived'] == 0], x='Age', fill=True, label='Not')
plt.title("Age Distribution of Survivors vs. Non-Survivors")
plt.legend()
plt.show()
```



general feature distribution as compared to survival

```
In [37]: # The impact of passenger class and gender on the survival outcome
pd.crosstab([df.Sex, df.Survived], df.Pclass, margins=True).style.background
```

Out[37]:

		Pclass	1	2	3	All
		Sex	Survived			
female	0	3	6	72	81	
	1	91	70	72	233	
male	0	77	91	300	468	
	1	45	17	47	109	
All		216	184	491	891	

We can observe from this that the likeliness of surviving as a female is significantly higher in Pclass 1 and 2, however in 3 the chance is 50/50

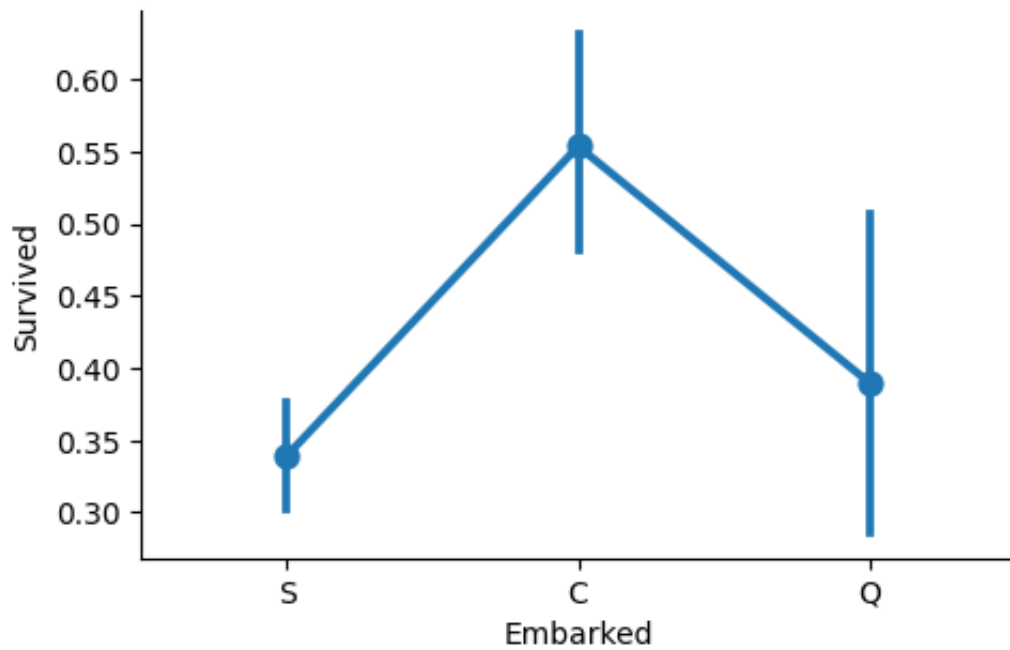
Whilst for males, the chance of survival is generally low across all the classes, we can observe a higher percentage of survival in Pclass 1 and 2 as compared to 3.

In [38]: *# Analysis of the embarkation point feature*
`pd.crosstab([df.Embarked, df.Pclass], [df.Sex, df.Survived], margins=True)`

Out[38]:

		Sex	female		male		All
		Survived	0	1	0	1	
Embarked	Pclass						
C	1	1	42	25	17	85	
	2	0	7	8	2	17	
	3	8	15	33	10	66	
Q	1	0	1	1	0	2	
	2	0	2	1	0	3	
	3	9	24	36	3	72	
S	1	2	48	51	28	129	
	2	6	61	82	15	164	
	3	55	33	231	34	353	
All		81	233	468	109	891	


In [39]: *# Use factorplot to visualize the relationship between the embarkation po*
`sns.catplot(x='Embarked', y='Survived', data=df, kind='point')`
`fig = plt.gcf()`
`fig.set_size_inches(5, 3)`
`plt.show()`



Evidently, the chances of survival is higher if you had embarked at C, however, this can be explained by the number of first class passengers embarking at C as compared to the other ports


```
In [40]: sns.countplot(x='Embarked', hue='Pclass', data=df)
fig.suptitle('Survival by Pclass, Sex, and Embarked')
```

```
Out[40]: Text(0.5, 0.98, 'Survival by Pclass, Sex, and Embarked')
```



As explained, the survival rate is directly correlated to class.

```
In [41]: # Visualize the relationship between age bands and survival, separated by
sns.catplot(x='Age_band', y='Survived', data=df, kind='point', col='Pclass')
plt.show()
```



We can see an overall decreasing trend in survivability with age, this can be explained by the "women and children first" procedure they had when allocating seats in liferafts

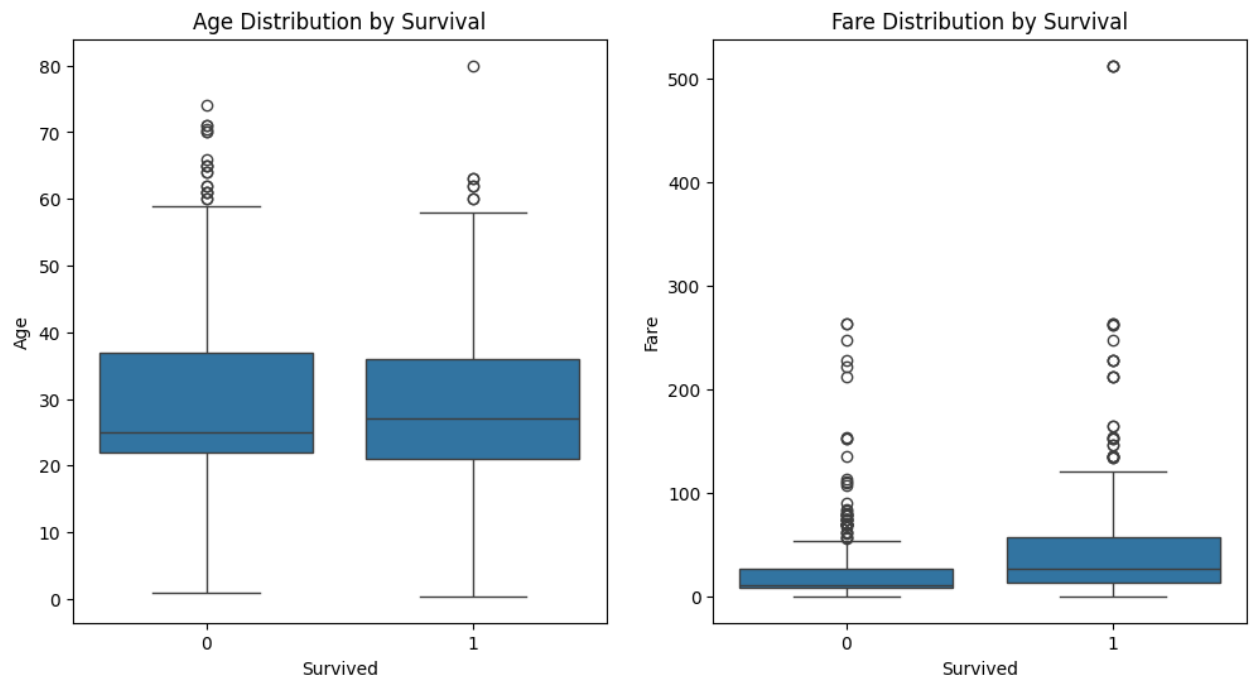
```
In [42]: # Fare and Age by Survival
fig, axes = plt.subplots(1, 2, figsize=(12, 6))
sns.boxplot(x='Survived', y='Age', data=df, ax=axes[0])
sns.boxplot(x='Survived', y='Fare', data=df, ax=axes[1])
```


```

axes[0].set_title("Age Distribution by Survival")
axes[1].set_title("Fare Distribution by Survival")
plt.show()

# Heatmap for Correlations
numeric_df = df.select_dtypes(include=['float64', 'int64'])
plt.figure(figsize=(10, 8))
sns.heatmap(numeric_df.corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Matrix of Titanic Dataset")
plt.show()

```





Hypothesis Testing

```
In [43]: # Chi-Square Test
contingency_table = pd.crosstab(df['Pclass'], df['Survived'])
chi2, p, dof, expected = chi2_contingency(contingency_table)
print(f"Chi-Square Test for Pclass and Survived: p-value = {p}")
```

Chi-Square Test for Pclass and Survived: p-value = 4.549251711298793e-23
 T-Test for Age and Survived: p-value = 0.07548530586360941


To test my first hypothesis that Pclass and survived is correlated, i decided to employ the chi square test, the result of the test indicated high correlation between survivability and class


```
In [44]: features = ['Age', 'Fare', 'Pclass', 'FamilySize']
X = df[features].fillna(0)
X_scaled = StandardScaler().fit_transform(X)

# PCA
pca = PCA(n_components=2)
```

```
X_pca = pca.fit_transform(X_scaled)
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=df['Survived'], cmap='viridis')
plt.title("PCA of Titanic Dataset")
plt.show()

# K-Means
kmeans = KMeans(n_clusters=3, random_state=0).fit(X_scaled)
df['Cluster'] = kmeans.labels_
sns.scatterplot(x=X_pca[:, 0], y=X_pca[:, 1], hue=df['Cluster'], palette=
plt.title("K-Means Clustering")
plt.show()
```





Model Introduction

```
In [45]: X = df[['Pclass', 'Sex', 'Age', 'Fare', 'FamilySize', 'IsAlone', 'HasCabin']]
X = pd.get_dummies(X, columns=['Sex'], drop_first=True)
y = df['Survived']

# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,

# Models
models = {
    'Logistic Regression': LogisticRegression(max_iter=1000),
    'Decision Tree': DecisionTreeClassifier(random_state=0),
    'Random Forest': RandomForestClassifier(random_state=0),
    'Gradient Boosting': GradientBoostingClassifier(random_state=0),
    'Support Vector Machine': SVC(probability=True, random_state=0),
    'Neural Network': MLPClassifier(max_iter=1000, random_state=0)
}

# Train and Evaluate
for model_name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    print(f"=== {model_name} ===")
    print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
```

```
print(f"Confusion Matrix:\n{confusion_matrix(y_test, y_pred)}\n")
```

=== Logistic Regression ===

Accuracy: 0.8060

Confusion Matrix:

```
[[144  24]
 [ 28  72]]
```

=== Decision Tree ===

Accuracy: 0.8134

Confusion Matrix:

```
[[146  22]
 [ 28  72]]
```

=== Random Forest ===

Accuracy: 0.8358

Confusion Matrix:

```
[[151  17]
 [ 27  73]]
```

=== Gradient Boosting ===

Accuracy: 0.8209

Confusion Matrix:

```
[[150  18]
 [ 30  70]]
```

=== Support Vector Machine ===

Accuracy: 0.7164

Confusion Matrix:

```
[[160   8]
 [ 68  32]]
```

=== Neural Network ===

Accuracy: 0.7910

Confusion Matrix:

```
[[142  26]
 [ 30  70]]
```